

```
In [3]: import pandas as pd
import numpy as np
import re
import string
from sklearn.model_selection import train_test_split
from sklearn.feature_extraction.text import CountVectorizer, TfidfVectorizer
from sklearn.naive_bayes import MultinomialNB
from sklearn.linear_model import LogisticRegression
from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer, WordNetLemmatizer
nltk.download('stopwords')
nltk.download('wordnet')
```

```
[nltk_data] Downloading package stopwords to C:\Users\GEETHA
[nltk_data]     SRI\AppData\Roaming\nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to C:\Users\GEETHA
[nltk_data]     SRI\AppData\Roaming\nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

Out[3]: True

```
In [4]: df = pd.read_csv('disaster_tweets_data(DS).csv')
print(df.head())
```

| | tweets | target |
|---|---|--------|
| 0 | Our Deeds are the Reason of this #earthquake M... | 1 |
| 1 | Forest fire near La Ronge Sask. Canada | 1 |
| 2 | All residents asked to 'shelter in place' are ... | 1 |
| 3 | 13,000 people receive #wildfires evacuation or... | 1 |
| 4 | Just got sent this photo from Ruby #Alaska as ... | 1 |

```
In [5]: df.isnull().sum()
df.dropna(inplace=True)
```

```
In [9]: def preprocess_text(text):
    text = text.lower() # Lowercase
    text = re.sub(r'@\w+', '', text) # Remove handles
    text = re.sub(r'http\S+|www\S+|https\S+', '', text) # Remove Links
    text = text.translate(str.maketrans('', '', string.punctuation)) # Remove punctuation
    tokens = text.split()
    stop_words = set(stopwords.words('english'))
    tokens = [word for word in tokens if word not in stop_words]
    lemmatizer = WordNetLemmatizer()
    tokens = [lemmatizer.lemmatize(word) for word in tokens] # Lemmatization
    return ' '.join(tokens)

df['cleaned_tweets'] = df['tweets'].apply(preprocess_text)
```

```
In [11]: # Choose one: CountVectorizer or TfidfVectorizer
vectorizer = TfidfVectorizer()
```

```
X = vectorizer.fit_transform(df['cleaned_tweets'])
y = df['target']
```

```
In [13]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_st
```

```
In [15]: # a) Multinomial Naive Bayes
```

```
nb = MultinomialNB()
nb.fit(X_train, y_train)
nb_pred = nb.predict(X_test)
```

```
# b) Logistic Regression
```

```
lr = LogisticRegression()
lr.fit(X_train, y_train)
lr_pred = lr.predict(X_test)
```

```
# c) KNN
```

```
knn = KNeighborsClassifier()
knn.fit(X_train, y_train)
knn_pred = knn.predict(X_test)
```

```
In [17]: models = {'Naive Bayes': nb_pred, 'Logistic Regression': lr_pred, 'KNN': knn_pred}
```

```
for name, pred in models.items():
    print(f"\n{name} Results:")
    print("Confusion Matrix:\n", confusion_matrix(y_test, pred))
    print("Classification Report:\n", classification_report(y_test, pred))
    print("Accuracy:", accuracy_score(y_test, pred))
```

Naive Bayes Results:

Confusion Matrix:

```
[[786 88]
 [217 432]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.78 | 0.90 | 0.84 | 874 |
| 1 | 0.83 | 0.67 | 0.74 | 649 |
| accuracy | | | 0.80 | 1523 |
| macro avg | 0.81 | 0.78 | 0.79 | 1523 |
| weighted avg | 0.80 | 0.80 | 0.80 | 1523 |

Accuracy: 0.7997373604727511

Logistic Regression Results:

Confusion Matrix:

```
[[790 84]
 [233 416]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.77 | 0.90 | 0.83 | 874 |
| 1 | 0.83 | 0.64 | 0.72 | 649 |
| accuracy | | | 0.79 | 1523 |
| macro avg | 0.80 | 0.77 | 0.78 | 1523 |
| weighted avg | 0.80 | 0.79 | 0.79 | 1523 |

Accuracy: 0.7918581746552856

KNN Results:

Confusion Matrix:

```
[[745 129]
 [226 423]]
```

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.77 | 0.85 | 0.81 | 874 |
| 1 | 0.77 | 0.65 | 0.70 | 649 |
| accuracy | | | 0.77 | 1523 |
| macro avg | 0.77 | 0.75 | 0.76 | 1523 |
| weighted avg | 0.77 | 0.77 | 0.76 | 1523 |

Accuracy: 0.7669074195666448

In []: